

```
In [1]: import numpy as np
import tensorflow
from matplotlib import pyplot as plt
from tensorflow.keras.models import Sequential, load_model
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, BatchNormalization
from tensorflow.keras.callbacks import ReduceLROnPlateau, EarlyStopping
from tensorflow.keras.regularizers import l2
from tensorflow.keras.datasets import cifar100
```

```
In [2]: (trainX, trainY), (testX, testY) = cifar100.load_data()
```

```
In [3]: # Target classes: numbers to text
# Source: https://github.com/keras-team/keras/issues/2653#issuecomment-450133996
classes = [
    'apple',
    'aquarium_fish',
    'baby',
    'bear',
    'beaver',
    'bed',
    'bee',
    'beetle',
    'bicycle',
    'bottle',
    'bowl',
    'boy',
    'bridge',
    'bus',
    'butterfly',
    'camel',
    'can',
    'castle',
    'caterpillar',
    'cattle',
    'chair',
    'chimpanzee',
    'clock',
    'cloud',
    'cockroach',
    'couch',
    'crab',
    'crocodile',
    'cup',
    'dinosaur',
    'dolphin',
    'elephant',
    'flatfish',
    'forest',
    'fox',
    'girl',
    'hamster',
    'house',
    'kangaroo',
    'computer_keyboard',
    'lamp',
    'lawn_mower',
    'leopard',
    'lion',
    'lizard',
    'lobster',
    'man',
    'maple_tree',
    'motorcycle',
    'mountain',
    'mouse',
    'mushroom',
```

```
'oak_tree',  
'orange',  
'orchid',  
'otter',  
'palm_tree',  
'pear',  
'pickup_truck',  
'pine_tree',  
'plain',  
'plate',  
'poppy',  
'porcupine',  
'possum',  
'rabbit',  
'raccoon',  
'ray',  
'road',  
'rocket',  
'rose',  
'sea',  
'seal',  
'shark',  
'shrew',  
'skunk',  
'skyscraper',  
'snail',  
'snake',  
'spider',  
'squirrel',  
'streetcar',  
'sunflower',  
'sweet_pepper',  
'table',  
'tank',  
'telephone',  
'television',  
'tiger',  
'tractor',  
'train',  
'trout',  
'tulip',  
'turtle',  
'wardrobe',  
'whale',  
'willow_tree',  
'wolf',  
'woman',  
'worm',  
]
```

```
In [4]: # Visualize first 9 images
for i in range(9):
    plt.subplot(330 + 1 + i)
    image = trainX[i]
    target = trainY[i][0]
    plt.axis('off')
    plt.imshow(image)
    plt.title(f'{classes[target]}')
plt.show()
```



```
In [5]: trainX = trainX.reshape(trainX.shape[0], 32,32,3)
trainY = trainY.reshape(trainY.shape[0], 1)
testX = testX.reshape(testY.shape[0], 32, 32, 3)
testY = testY.reshape(testY.shape[0], 1)
```

```
In [6]: # Parse numbers as floats
trainX = trainX.astype('float32')
testX = testX.astype('float32')

# Normalize data
trainX = trainX / 255.0
testX = testX / 255.0
```

```
In [7]: model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape = (32, 32, 3), kernel_regularizer = l
2(0.005), activation = 'elu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size = (2, 2)))
model.add(Conv2D(128, (3, 3), kernel_regularizer = l2(0.005), activation = 'el
u'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size = (2, 2)))
model.add(Conv2D(128, (3, 3), kernel_regularizer = l2(0.005), activation = 'el
u'))
model.add(BatchNormalization())
model.add(Conv2D(256, (3, 3), kernel_regularizer = l2(0.005), activation = 'el
u'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size = (2, 2)))
model.add(Flatten())
model.add(Dense(units = 128, activation = 'elu'))
model.add(Dense(units = 256, activation = 'elu'))
model.add(Dense(units = 100, activation = 'softmax'))
```

```
In [8]: model.compile(optimizer = 'adam', loss = 'sparse_categorical_crossentropy', me
trics = ['accuracy'])
```

```
In [9]: rlrop = ReduceLROnPlateau(monitor = 'val_loss', factor = 0.1, patience = 10)
```

```
In [10]: early_stop = EarlyStopping(monitor = 'val_loss',
                                     mode = 'min',
                                     verbose = 1,
                                     patience = 30)
```

```
In [11]: history = model.fit(trainX, trainY,  
                             steps_per_epoch = 400,  
                             epochs = 100,  
                             #verbosity = 1,  
                             validation_split = 0.2,  
                             callbacks = [rlrop, early_stop])
```

Train on 40000 samples, validate on 10000 samples

```
Epoch 1/100
40000/40000 [=====] - 126s 3ms/sample - loss: 4.3252
- accuracy: 0.1828 - val_loss: 5.2625 - val_accuracy: 0.0700
Epoch 2/100
40000/40000 [=====] - 104s 3ms/sample - loss: 3.2391
- accuracy: 0.2998 - val_loss: 4.1788 - val_accuracy: 0.1761
Epoch 3/100
40000/40000 [=====] - 104s 3ms/sample - loss: 3.0300
- accuracy: 0.3465 - val_loss: 4.3409 - val_accuracy: 0.1775
Epoch 4/100
40000/40000 [=====] - 106s 3ms/sample - loss: 2.9066
- accuracy: 0.3832 - val_loss: 3.3714 - val_accuracy: 0.2982
Epoch 5/100
40000/40000 [=====] - 107s 3ms/sample - loss: 2.8178
- accuracy: 0.4107 - val_loss: 3.2840 - val_accuracy: 0.3373
Epoch 6/100
40000/40000 [=====] - 107s 3ms/sample - loss: 2.7530
- accuracy: 0.4304 - val_loss: 3.5828 - val_accuracy: 0.2937
Epoch 7/100
40000/40000 [=====] - 109s 3ms/sample - loss: 2.6902
- accuracy: 0.4507 - val_loss: 3.6019 - val_accuracy: 0.2947
Epoch 8/100
40000/40000 [=====] - 110s 3ms/sample - loss: 2.6293
- accuracy: 0.4685 - val_loss: 3.5475 - val_accuracy: 0.3119
Epoch 9/100
40000/40000 [=====] - 108s 3ms/sample - loss: 2.5816
- accuracy: 0.4820 - val_loss: 3.2323 - val_accuracy: 0.3699
Epoch 10/100
40000/40000 [=====] - 109s 3ms/sample - loss: 2.5296
- accuracy: 0.4986 - val_loss: 3.4871 - val_accuracy: 0.3417
Epoch 11/100
40000/40000 [=====] - 116s 3ms/sample - loss: 2.4882
- accuracy: 0.5110 - val_loss: 3.4263 - val_accuracy: 0.3437
Epoch 12/100
40000/40000 [=====] - 110s 3ms/sample - loss: 2.4410
- accuracy: 0.5217 - val_loss: 3.5240 - val_accuracy: 0.3458
Epoch 13/100
40000/40000 [=====] - 110s 3ms/sample - loss: 2.3977
- accuracy: 0.5376 - val_loss: 3.7135 - val_accuracy: 0.3152
Epoch 14/100
40000/40000 [=====] - 110s 3ms/sample - loss: 2.3695
- accuracy: 0.5473 - val_loss: 3.5284 - val_accuracy: 0.3522
Epoch 15/100
40000/40000 [=====] - 110s 3ms/sample - loss: 2.3311
- accuracy: 0.5572 - val_loss: 4.1948 - val_accuracy: 0.2615
Epoch 16/100
40000/40000 [=====] - 110s 3ms/sample - loss: 2.2837
- accuracy: 0.5714 - val_loss: 3.6979 - val_accuracy: 0.3291
Epoch 17/100
40000/40000 [=====] - 108s 3ms/sample - loss: 2.2457
- accuracy: 0.5809 - val_loss: 3.6579 - val_accuracy: 0.3471
Epoch 18/100
40000/40000 [=====] - 109s 3ms/sample - loss: 2.2117
- accuracy: 0.5953 - val_loss: 3.4908 - val_accuracy: 0.3743
Epoch 19/100
40000/40000 [=====] - 109s 3ms/sample - loss: 2.1848
```

```
- accuracy: 0.6023 - val_loss: 3.8797 - val_accuracy: 0.3349
Epoch 20/100
40000/40000 [=====] - 109s 3ms/sample - loss: 1.6245
- accuracy: 0.7621 - val_loss: 2.8346 - val_accuracy: 0.4920
Epoch 21/100
40000/40000 [=====] - 110s 3ms/sample - loss: 1.3456
- accuracy: 0.8293 - val_loss: 2.7888 - val_accuracy: 0.4940
Epoch 22/100
40000/40000 [=====] - 110s 3ms/sample - loss: 1.1831
- accuracy: 0.8631 - val_loss: 2.7597 - val_accuracy: 0.4916
Epoch 23/100
40000/40000 [=====] - 109s 3ms/sample - loss: 1.0601
- accuracy: 0.8828 - val_loss: 2.7660 - val_accuracy: 0.4902
Epoch 24/100
40000/40000 [=====] - 109s 3ms/sample - loss: 0.9537
- accuracy: 0.9054 - val_loss: 2.7583 - val_accuracy: 0.4861
Epoch 25/100
40000/40000 [=====] - 110s 3ms/sample - loss: 0.8592
- accuracy: 0.9222 - val_loss: 2.8006 - val_accuracy: 0.4859
Epoch 26/100
40000/40000 [=====] - 110s 3ms/sample - loss: 0.7780
- accuracy: 0.9370 - val_loss: 2.8942 - val_accuracy: 0.4713
Epoch 27/100
40000/40000 [=====] - 109s 3ms/sample - loss: 0.7146
- accuracy: 0.9471 - val_loss: 2.9108 - val_accuracy: 0.4803
Epoch 28/100
40000/40000 [=====] - 109s 3ms/sample - loss: 0.6554
- accuracy: 0.9561 - val_loss: 2.9620 - val_accuracy: 0.4747
Epoch 29/100
40000/40000 [=====] - 110s 3ms/sample - loss: 0.6076
- accuracy: 0.9629 - val_loss: 2.9858 - val_accuracy: 0.4788
Epoch 30/100
40000/40000 [=====] - 108s 3ms/sample - loss: 0.5648
- accuracy: 0.9691 - val_loss: 3.0552 - val_accuracy: 0.4682
Epoch 31/100
40000/40000 [=====] - 109s 3ms/sample - loss: 0.5256
- accuracy: 0.9735 - val_loss: 3.0819 - val_accuracy: 0.4673
Epoch 32/100
40000/40000 [=====] - 110s 3ms/sample - loss: 0.4995
- accuracy: 0.9735 - val_loss: 3.1257 - val_accuracy: 0.4728
Epoch 33/100
40000/40000 [=====] - 110s 3ms/sample - loss: 0.4754
- accuracy: 0.9768 - val_loss: 3.1471 - val_accuracy: 0.4680
Epoch 34/100
40000/40000 [=====] - 109s 3ms/sample - loss: 0.4442
- accuracy: 0.9816 - val_loss: 3.2416 - val_accuracy: 0.4683
Epoch 35/100
40000/40000 [=====] - 110s 3ms/sample - loss: 0.3966
- accuracy: 0.9916 - val_loss: 3.1433 - val_accuracy: 0.4812
Epoch 36/100
40000/40000 [=====] - 110s 3ms/sample - loss: 0.3803
- accuracy: 0.9952 - val_loss: 3.1494 - val_accuracy: 0.4813
Epoch 37/100
40000/40000 [=====] - 110s 3ms/sample - loss: 0.3728
- accuracy: 0.9965 - val_loss: 3.1493 - val_accuracy: 0.4798
Epoch 38/100
40000/40000 [=====] - 108s 3ms/sample - loss: 0.3674
```



```
- accuracy: 0.9967 - val_loss: 3.1642 - val_accuracy: 0.4828
Epoch 39/100
40000/40000 [=====] - 108s 3ms/sample - loss: 0.3614
- accuracy: 0.9972 - val_loss: 3.1652 - val_accuracy: 0.4828
Epoch 40/100
40000/40000 [=====] - 108s 3ms/sample - loss: 0.3565
- accuracy: 0.9970 - val_loss: 3.1896 - val_accuracy: 0.4803
Epoch 41/100
40000/40000 [=====] - 108s 3ms/sample - loss: 0.3501
- accuracy: 0.9978 - val_loss: 3.1823 - val_accuracy: 0.4812
Epoch 42/100
40000/40000 [=====] - 108s 3ms/sample - loss: 0.3453
- accuracy: 0.9975 - val_loss: 3.1954 - val_accuracy: 0.4819
Epoch 43/100
40000/40000 [=====] - 109s 3ms/sample - loss: 0.3405
- accuracy: 0.9979 - val_loss: 3.1892 - val_accuracy: 0.4806
Epoch 44/100
40000/40000 [=====] - 110s 3ms/sample - loss: 0.3357
- accuracy: 0.9980 - val_loss: 3.2206 - val_accuracy: 0.4770
Epoch 45/100
40000/40000 [=====] - 108s 3ms/sample - loss: 0.3297
- accuracy: 0.9985 - val_loss: 3.2055 - val_accuracy: 0.4814
Epoch 46/100
40000/40000 [=====] - 109s 3ms/sample - loss: 0.3293
- accuracy: 0.9988 - val_loss: 3.2074 - val_accuracy: 0.4809
Epoch 47/100
40000/40000 [=====] - 112s 3ms/sample - loss: 0.3288
- accuracy: 0.9987 - val_loss: 3.2057 - val_accuracy: 0.4813
Epoch 48/100
40000/40000 [=====] - 110s 3ms/sample - loss: 0.3279
- accuracy: 0.9987 - val_loss: 3.2077 - val_accuracy: 0.4806
Epoch 49/100
40000/40000 [=====] - 110s 3ms/sample - loss: 0.3263
- accuracy: 0.9987 - val_loss: 3.2090 - val_accuracy: 0.4803
Epoch 50/100
40000/40000 [=====] - 108s 3ms/sample - loss: 0.3267
- accuracy: 0.9987 - val_loss: 3.2099 - val_accuracy: 0.4807
Epoch 51/100
40000/40000 [=====] - 108s 3ms/sample - loss: 0.3255
- accuracy: 0.9987 - val_loss: 3.2109 - val_accuracy: 0.4812
Epoch 52/100
40000/40000 [=====] - 109s 3ms/sample - loss: 0.3251
- accuracy: 0.9985 - val_loss: 3.2109 - val_accuracy: 0.4821
Epoch 53/100
40000/40000 [=====] - 109s 3ms/sample - loss: 0.3247
- accuracy: 0.9986 - val_loss: 3.2120 - val_accuracy: 0.4811
Epoch 54/100
40000/40000 [=====] - 108s 3ms/sample - loss: 0.3242
- accuracy: 0.9986 - val_loss: 3.2163 - val_accuracy: 0.4820
Epoch 00054: early stopping
```

```
In [12]: model.save('Version_Final.h5')
```

In [13]: `model.summary()`

Model: "sequential"


Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 30, 30, 32)	896
batch_normalization (Batch Normalization)	(None, 30, 30, 32)	128
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 128)	36992
batch_normalization_1 (Batch Normalization)	(None, 13, 13, 128)	512
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 128)	0
conv2d_2 (Conv2D)	(None, 4, 4, 128)	147584
batch_normalization_2 (Batch Normalization)	(None, 4, 4, 128)	512
conv2d_3 (Conv2D)	(None, 2, 2, 256)	295168
batch_normalization_3 (Batch Normalization)	(None, 2, 2, 256)	1024
max_pooling2d_2 (MaxPooling2D)	(None, 1, 1, 256)	0
flatten (Flatten)	(None, 256)	0
dense (Dense)	(None, 128)	32896
dense_1 (Dense)	(None, 256)	33024
dense_2 (Dense)	(None, 100)	25700
=====		
Total params: 574,436		
Trainable params: 573,348		
Non-trainable params: 1,088		

In [14]: `# Generate generalization metrics`  
`score = model.evaluate(testX, testY, verbose=0)`  
`print(f'Test loss: {score[0]} / Test accuracy: {score[1]}')`


Test loss: 3.152114510726929 / Test accuracy: 0.4839000105857849

In [15]: `testY_pred = model.predict(testX)`

```
In [16]: # Visualize history
# Plot history: Loss
plt.plot(history.history['val_loss'], label = 'val_loss')
plt.plot(history.history['loss'], label = 'train_loss')
plt.title('Validation loss history')
plt.ylabel('Loss value')
plt.xlabel('No. epoch')
plt.legend()
plt.show()
```




```
In [17]: # Plot history: Accuracy
plt.plot(history.history['val_accuracy'], label = 'validation accuracy')
plt.plot(history.history['accuracy'], label = 'train accuracy')
plt.title('Validation accuracy history')
plt.ylabel('Accuracy value (%)')
plt.xlabel('No. epoch')
plt.legend()
plt.show()
```



```
In [18]: import random
fig=plt.figure(figsize=(10,7))
columns = 3
rows = 3
# Visualize first 9 images
for i in range(1,10):
    fig.add_subplot(rows, columns, i)
    k = random.randrange(10000)
    image = testX[k]
    target_pred = testY_pred[k].tolist().index(testY_pred[k].max())
    target_actual = testY[k][0]
    plt.axis('off')
    plt.imshow(image)
    plt.title(classes[target_pred]+' ('+classes[target_actual]+''))
plt.show()
```


plain (plain)



motorcycle (motorcycle)



sunflower (sunflower)



rabbit (ray)



caterpillar (caterpillar)



chair (chair)



leopard (leopard)



fox (tiger)



plain (road)

